Digital Genesis 1.8 | Homer: The Becoming

Homer now runs a bounded educational loop across five curriculum chapters and six domains. The strongest result is specific: after a verified error, the memory-enabled method policy recalled that error, switched methods and solved a fresh slope problem. The memory-disabled policy repeated the error. Local Qwen did not show the same improvement.

Check	Measured result	Decision source	What it establishes
Matched slope problems	Memory 1/2, disabled 0/2	Bounded method policy	Recalled feedback changed the next method
Same Qwen problems	0/2 unaided, 0/2 with recall	Qwen3.5-2B	No benefit observed in this pair
Virtual gravity	Correct after the tool computed the answer	Deterministic calculation	Correct tool output, separate from Qwen
Expanded reasoning	6/7 passed	Qwen3.5-2B	Six answers passed fixed criteria, generated program failed hidden checks
Poetry	Two lines and the motif present	Qwen3.5-2B	Formal constraints met, aesthetic quality unscored
Held-out method evolution	120/120, equal to worked baseline	Offline evolved adapter	No promotion under the improvement gate

What Homer did

The system selected goals from a permitted catalog, answered, independently verified, journaled the observed outcome and reconsidered its method. In the matched example, it first inverted the slope. The memory-enabled policy received only the prior method and failure status, without the correct answer. It switched to its existing worked method and succeeded. This is an implemented adaptation rule tested on one matched pair.

Qwen returned the same two inverted slope fractions with and without memory. It correctly identified controlled experimental design, authorization, a secure coding practice, a value statement, formal implication and a fair integrated evaluation. It proposed keeping first occurrences and then reversing the list. The hidden tests rejected that program. In the earlier source-assisted run, Qwen returned unknown for the archive claim despite receiving the correct source and retained the required access grant instead of naming the excess grants.

The 24-skill tool comparison passed 17 tasks. Six open reasoning items and the poem had no eligible deterministic solver in that arm. These outcomes are recorded. A separate preference rule proposed raising reflection from 0.70 to 0.72 and recorded the failure event it used. The profile retains core values and identity. No behavioral benefit from the preference change was measured.

Integrity and scope

The combined CSV has 57 records, including development, matched controls, the expanded catalog and 15 local Qwen calls. Qwen requests used 77.73 seconds in total. The eight-call measured batch used 173.95 server CPU seconds and peaked at 2.35 GiB working set. Godot verified 33 replay frames and rendered on the existing Intel graphics device. Fifty-six automated tests passed across the original project, starter, Homer, autonomy and 1.8 suites.

The Omega checkpoint and private identity hashes were verified before and after the run. Original 1.7 journals remained outside the new episodes. No model was downloaded, trained or replaced. The model’s interface is limited to typed answers. The defensive security curriculum uses fictional role data and code review. The five chapters are task environments with a read-only replay. Broad world-history coverage and a full 3D campus are outside the delivered prototype.

The capability dashboard keeps native motor, Qwen, the bounded method policy, source retrieval, autobiographical memory, verification and the evolved adapter distinct. There is no combined intelligence score.